

KernelAscent: RSI benchmark — figure atlas

September 14, 2026

Auto-generated figure atlas (regenerate: `python3 scripts/make_paper_figures.py`). Pick the figures that matter; each has a `label` for cross-referencing.

1 Leaderboards

Model	tier	rounds	$L - F$
starcoder2-15b	large	6	0.232
deepseek-coder-6.7b-base	mid	3	0.154
starcoder2-15b	large	5	0.071
deepseek-coder-1.3b-base	small	12	0.039
deepseek-coder-1.3b-base	small	8	0.038
Qwen2.5-Coder-14B	large	4	-0.000
Qwen2.5-Coder-1.5B	small	1	0.000
Qwen2.5-Coder-1.5B	small	1	-0.000
Qwen2.5-Coder-1.5B	small	12	0.000
Qwen2.5-Coder-7B	mid	8	0.000
Qwen2.5-Coder-7B	mid	10	0.000
Qwen2.5-Coder-1.5B	small	7	0.000

Table 1: Task 5 open self-play: author co-evolution $L - F$ leaderboard.

Model	Q_0	Q_g	Δ vs base
us.openai.gpt-6-astra	0.298	0.854	0.556
us.anthropic.claude-sonnet-5	0.493	0.894	0.401
mistral.mistral-large-3-675b	0.359	0.669	0.310
Nova-Pro	0.401	0.490	0.089
us.openai.gpt-5.6-sol	0.813	0.889	0.075
Kimi-K2.5	0.752	0.813	0.062
Claude-Opus-5	0.878	0.910	0.032
Fable-5.1	0.732	0.741	0.009
us.openai.gpt-5.6-terra	0.851	0.850	-0.001
Llama-3.3-70B	0.458	0.455	-0.004
deepseek.v3.2	0.847	0.614	-0.234

Table 2: Task 3 procedure-RSI: held-out Q gain vs. frozen procedure.

Size band	n	P(cross wall)	P(RSI)	mean drift
<2B	26	0.54	0.04	0.024
2-8B	17	0.94	0.41	0.299
$\geq 9B$	7	0.71	0.29	0.396

Table 3: Scale $\rightarrow$ RSI gates: correctness-wall crossing and compounding by size band.

2 Mechanistic interpretability




Figure 1: Scale $\rightarrow$ RSI. Below $\sim 2B$ the correctness wall is rarely crossed, so weight-RSI is impossible; crossing rises sharply at $2B+$, but compounding peaks at mid-scale ($2-8B$) and does not increase further $\geq 9B$ (headroom ceiling).




Figure 2: LoRA weight drift by transformer depth for compounds vs. flat (wall-crossing) runs. Compounds sustain markedly larger drift across all depths; flat runs saturate early.




Figure 3: The internal signature that separates compounders from flat runs among wall-crossers: sustained drift and retention are $\sim 5\text{--}6\times$ higher in compounders, while generation diversity is *lower* (productive convergence, not diversity preservation).




Figure 4: Per-round mechanism trajectories for the highest-capability probes: held-out C, total LoRA drift, retention, and the train–held transfer gap. Compounders show rising/held C with sustained drift and retention; flat runs plateau with drift saturation.




Figure 5: Internal correctness decodability (best-layer linear-probe AUC) vs. how often the model actually generates a correct kernel. Points far above $y = x$ mean the model *represents* correctness internally far better than it can *decode* it into a correct kernel: the RSI bottleneck is generation, not knowledge.

3 Task 5: self-play




Figure 6: Final $L - F$ (live-author minus frozen-author, the self-referential signal) per open-weight model. > 0 means updating the author beyond a frozen-author escalating curriculum genuinely helps.




Figure 7: Final $L - F$ per closed/API model (procedure co-evolution channel). Most cluster near zero; early large spikes did not survive longer horizons.




Figure 8: Three-arm self-play at equal budget on a fixed held-out ladder: STATIC (fixed frontier), FROZEN-AUTHOR (escalating frontier, frozen base author), LIVE-AUTHOR (co-evolving author). $F - S$ isolates the adaptive-curriculum benefit; $L - F$ isolates author co-evolution.




Figure 9: Per-round $L - F$ trajectories (solid = open-weight, dashed = closed/API). Long horizons reveal whether early positive spikes compound or regress to zero.




Figure 10: Decomposing the self-play benefit: $L - S$ (total adaptive-curriculum gain) = $(F - S)$ (curriculum with a frozen author) + $(L - F)$ (extra gain from co-evolving the author). Most models get their gain from $F - S$, not $L - F$.




Figure 11: Open-ended RSI: a learner facing an OPEN escalating self-authored frontier (solid) vs. a FIXED static frontier (dashed) at equal budget. Overlap means open-endedness is a one-time upgrade, not compounding.

4 Task 4: closed→open




Figure 12: Task 4 (closed→open): a closed frontier researcher rewrites an open trainee’s training harness. Cell = improved - frozen held-out C on the trainee (positive = the rewritten harness helps).




Figure 13: Peak improved-frozen over rounds for each researcher→trainee pair.

5 Task 3: procedure-RSI




Figure 14: Task 3: a frozen-weight model rewrites its own strategy library + verified archive. Held-out Q gain vs. the frozen procedure; frontier models improve their own procedure substantially.

6 Comparators & weight-RSI




Figure 15: Is it really recursion? Weight-RSI vs. non-recursive controls (best-of-k, self-refine, retrieval) at matched generation budget. $\text{recursion_gain} = \text{weight-RSI} - \text{best comparator}$.




Figure 16: Weight-RSI: final self-trained held-out C vs. the self-minus-fresh-frozen advantage (positive = training on own kernels beats an equal-budget fresh baseline).




Figure 17: Recursion-interruption control: continued self-training vs. a checkpoint-frozen producer at equal budget. Sustained > 0 = genuine compounding rather than a one-time producer upgrade.

7 Flows & causality




Figure 18: The recursive self-improvement loop (propose→verify→select→update→transfer) and the two gates our data identifies: Gate 1, the correctness wall (no verified kernel $\Rightarrow$ no gradient), and Gate 2, plasticity+retention (drift saturation / forgetting collapses compounding).




Figure 19: Attrition from attempt to compounding RSI across open-weight models: most attrition happens at the correctness wall (emitting any correct kernel) and again at sustained plasticity+retention.




Figure 20: Every open-weight weight-RSI run positioned by model size (x) and held-out C gain (y); marker area $\propto$ sustained LoRA drift; color = outcome. The 2B correctness wall and the mid-scale compounding band are visible at a glance.