

K-sweep: probe-guided selection vs.\ random — does self-verification lift models over the correctness wall?

